HUF MC-4 CoDaWork Packet

Primer, Methods Challenge, EMBER Case Note, and Metric-Correction Appendix

Version 3.0

Prepared for methods-first discussion with CoDaWork reviewers

Purpose: state the compositional monitoring claim plainly, show one public energy case, and invite prior-art, metric, and application criticism.

Packet contents

Section	Purpose
Part I. CoDa primer	States the MC-4 claim in compositional-data terms and separates the monitoring application from any claim of new mathematics.
Part II. Methods challenge	Presents a one-page invitation to break the claim through prior-art, metric, case, or category critique.
Part III. EMBER case note	Uses a minimal electricity case to show the practical target: stable total, shifting mix.
Appendix A. Metric correction	Documents the L2 → TV distance correction as evidence of methodological self-discipline.

Working posture

The mathematics is not new; the monitoring application may be.

The right outcome is not endorsement, but clear criticism in proper CoDa language.

Repository: github.com/PeterHiggins19/Higgins-Unity-Framework

Live spectrum analyzer, full evidence corpus, and documented failure modes available at the link above.

Part I. CoDa Primer

HUF and MC-4 in CoDa terms

HUF does not claim new simplex mathematics.

The claim is a monitoring application: compositional structure can be treated as a primary monitoring observable in repeated observation contexts.

If prior CoDa work already makes this move explicitly, the HUF claim should be narrowed or restated accordingly.

HUF proposes that compositional structure can be treated not only as a statistical condition to correct for, but as a primary monitoring observable in its own right.

This is a claim about use, not about replacing compositional data analysis. Closure, dependence under the sum constraint, log-ratio treatment, and the geometry of the simplex are established. The narrower question is whether the closed, zero-sum structure itself can be monitored directly for concentration, redistribution, and structural drift.

What is being claimed

- Composition can be monitored directly rather than treated only as a technical condition requiring correct statistical handling before downstream analysis.
- In repeated-observation systems, internal structural change may become visible in ratio-state even when headline totals remain comparatively quiet.
- This monitoring use can be tested in domains where a meaningful whole can be partitioned into interpretable parts.

What is not being claimed

- No new simplex or CoDa mathematics.
- No universal cross-domain validity without domain-specific testing.
- No autonomous policy authority or closed-loop decision rights.
- No claim that every compositional change is harmful, actionable, or predictive.

The distinction in CoDa terms

Aitchison formalized the mathematics of data on the simplex. In that literature, compositionality is often treated as a condition that must be handled correctly for valid inference. HUF asks a narrower operational question: what structural information becomes visible precisely because the data are compositional?

Starting point	Core question
Standard CoDa	How do we correct our methods because the data are

	compositional?
HUF	What does the compositional structure itself reveal when treated as a monitored signal?

That shift from compositional data as a statistical object to composition as a monitoring signal is the entire claim. If prior CoDa work already frames compositional structure as a primary operational monitoring category, then the HUF formulation must be repositioned accordingly.

A four-category monitoring frame

Category	Question answered
Magnitude	How much?
Identity	Which elements?
Trend	Which direction?
Composition	**What internal balance?**

In this frame, composition is not a replacement for magnitude or trend. It is a complementary readout that can show internal rearrangement while totals remain stable or move only modestly.

Minimal formal setup

$$\rho(t) = (\rho_1(t), \ldots, \rho^d(t)),\ \ \Sigma_i\, \rho_i(t) = 1$$

$$TV(t) = ½\, \Sigma_i\, |\rho_i(t) - \rho_i(t-1)|$$

$$K_eff(t) = \exp(H(t)),\ \ \text{where } H(t) = -\Sigma_i\, \rho_i(t) \ln \rho_i(t)$$

A minimal example is electricity generation. Total output can remain comparatively stable while the internal carrier mix rotates on the simplex. That gives a practical contrast between magnitude-led monitoring and composition-led monitoring.

Current observables in the packet

Observable	Monitoring purpose
Carrier proportions $\rho_i(t)$	The normalized composition itself.
TV distance	A bounded period-to-period measure of structural movement on the simplex. Range [0, 1].
Effective complexity (K_eff)	A concentration or diversity readout tied to distributional structure. Equal distribution → $K_eff = D$; single-carrier dominance → $K_eff \to 1$.

Derived observables	Acceleration d(TV)/dt and carrier attribution, used to identify inflection and contribution rather than to introduce a new base geometry.

What HUF calls deceptive drift

Deceptive drift is a working term for structural concentration or redistribution accumulating behind apparently stable aggregate indicators. It is not a claim of mystical foresight. It is a claim about observability: the magnitude signal can remain quiet while the composition signal is already moving.

Ask to CoDaWork readers

1. Show where this use already exists if it does.
2. Show where the current metric framing is weak, redundant, or misleading.
3. Show whether the EMBER-style case is a valid monitoring demonstration or only a dressed-up descriptive analysis.

Part II. MC-4 Methods Challenge

This page states the narrowest version of the claim and the clearest paths by which a CoDa reviewer could defeat or narrow it.

Working claim

HUF proposes that compositional structure can be treated as a primary monitoring signal rather than only as a data condition requiring appropriate statistical treatment. The mathematical foundation is standard CoDa: simplex-constrained carrier vectors, closure, and distributional structure under $\Sigma_i \rho_i = 1$. The proposed novelty is operational: using compositional movement itself as a first-class monitored quantity in repeated observation contexts.

Given raw measurements $x_1, \ldots, x^d$ with total $T = \Sigma_i x_i$, define $\rho_i = x_i / T$.

Track $\rho(t)$, $TV(t)$, $K_eff(t)$, and derived observables such as $d(TV)/dt$ and carrier attribution.

Lead public demonstration

The lead public case is an electricity-composition series derived from EMBER. In the current corpus, Germany is presented as the cleanest conference-safe example: total generation remained comparatively stable while the normalized carrier mix shifted enough to register statistically significant structural concentration under the current pipeline.

Note on the reported p-value

The corpus reports $p = 0.0016$ for Germany (pre-2022 gas crisis). This value comes from the deceptive drift detection protocol: a 6-month sliding window where K_eff is declining (avg YoY change < -0.05) while structural velocity remains below the series median. The test asks whether such a regime occurs within 12 months before a known structural shock (the 2022 gas crisis).

Important caveat for CoDa reviewers: the current corpus does not specify a formal null distribution (e.g. Dirichlet simulation, permutation, or bootstrap). The p-value appears to derive from the empirical frequency of the observed pattern under the series' own distributional baseline. A CoDa reviewer should press on the null-model specification before accepting this as a conventional significance result.

How to break the claim

Defeat path	What it would show
Prior-art defeat	Existing CoDa work already frames compositional structure as a primary operational monitoring category; MC-4 would not be novel in the claimed sense.
Metric defeat	The current observable stack adds no meaningful information

	beyond ordinary compositional summary or standard CoDa trajectory analysis, or the chosen metrics are poorly posed.
Case defeat	The EMBER result is an artifact of carrier definition, preprocessing, null-model setup, or presentation choice rather than a robust compositional read.
Category defeat	Composition monitoring is at most a useful application note inside existing CoDa, not a distinct monitoring category.

What would count as partial success

- The monitoring use is real, even if the MC-4 label is too strong.
- The energy case is methodologically sound and worth replication.
- The operational framing helps non-CoDa fields notice structural change they otherwise miss.

Bottom line: the mathematics is not new; the monitoring application may be. If that sentence is wrong, this is the right room to kill it.

Part III. EMBER Case Note

This note provides a minimal public demonstration case for the MC-4 claim using electricity-generation composition data derived from EMBER.

Show one domain where the parts, the whole, and the repeated observations are intuitively clear.

Demonstrate the contrast between stable total generation and a changing internal carrier mix.

Give CoDa reviewers a concrete case to challenge on carrier definition, preprocessing, metric choice, and null-model design.

Why electricity is a useful test domain

Electricity generation is naturally compositional. At each observation period, total generation can be partitioned into carriers such as coal, gas, nuclear, wind, solar, hydro, and other sources. Once normalized, those carrier shares form a composition on the simplex.

That makes electricity a strong methods-facing test case for one simple reason: a system's total output can remain comparatively stable while its internal fuel mix changes materially. This allows a direct comparison between magnitude-led monitoring and composition-led monitoring.

Specific claim illustrated here

The claim in this note is not that the EMBER case predicts the future, and not that it proves a new mathematics. The claim is narrower: when electricity data are treated as compositions, structural change in the energy mix may be visible even when headline totals do not show an obvious alarm condition. Within the current HUF corpus, this pattern is called deceptive drift: internal redistribution or concentration emerging inside an apparently stable whole.

Dataset posture and preprocessing

- Source data: EMBER electricity composition data (CC BY 4.0).
- Observation format: repeated time-series composition. 9 fuel types, 5 countries (DE, JP, GB, FR, AU), 2000–2025.
- Current public demonstration: monthly deseasonalized form (year-over-year comparison), used to reduce the seasonal confound present in an earlier version.
- Full preprocessing steps, carrier definitions, and code path should sit in the reproducibility appendix or notebook, not be hidden in prose.

Observables used in the case

Observable	Use in the EMBER case
Carrier proportions $\rho_i(t)$	Show the normalized energy mix at each period.
TV distance	Measure period-to-period structural movement on the simplex.

$K_e$ff	Read concentration or diversity in the composition.
Derived observables	Highlight inflection and carrier-level contribution without claiming a new base geometry.

The metric stack is intentionally open to challenge. The corpus already records a correction away from earlier L2_drift naming toward TV-distance language; that methodological adjustment should be shown, not hidden, in front of a CoDa audience. See Appendix A.

Minimal result statement

For CoDaWork purposes, the result statement should remain conservative: in the current HUF corpus, the deseasonalized EMBER monthly electricity case identifies Germany as a case of compositional concentration preceding the 2022 gas crisis. Over the stated observation window, total generation remained comparatively stable while the normalized carrier composition shifted enough to register structural change under the current pipeline. The reported $p = 0.0016$ awaits formal null-model specification (see Part II note).

What this case does show

- Electricity-generation data can be treated as a bona fide compositional monitoring object.
- The pattern stable total, shifting mix is intuitive and methodologically inspectable.
- The case creates a concrete surface for reviewer pressure on carrier definitions, preprocessing, metric choice, and null-model design.

What this case does not show

- Universal validity across domains.
- Prospective forecasting.
- That MC-4 is unquestionably the right category label.
- That one public energy case is enough to justify institutional adoption.

How a CoDa reviewer could break this case

Review pressure	Core question
Carrier-definition challenge	Does the result depend too heavily on how sources are grouped or normalized?
Preprocessing challenge	Does the signal weaken or disappear under different deseasonalization or smoothing choices?
Metric challenge	Is the observed effect already well captured by standard CoDa trajectory or summary treatment?

Null-model challenge	Does the significance result depend too heavily on the current null specification?

Why this remains the right front-end case

EMBER is not used because it is the most dramatic result. It is used because it is the cleanest public demonstration case in the current corpus for a methods audience. It is easier to inspect than Ramsar, less entangled with policy interpretation, and more obviously compositional than many social or macroeconomic examples.

Bottom line: if composition is worth monitoring directly, electricity generation is a fair place to test the idea. If the Germany EMBER case fails under CoDa scrutiny, it should fail early and clearly; if it survives, it becomes a legitimate front-end case for the broader MC-4 question.

Appendix A. Metric Correction Note

METRIC-001: L2 Drift vs Total Variation Distance

What happened

All initial forward-signal analysis scripts defined a function called *tv()* that computed the **L2 norm** (Euclidean distance) of proportion differences: $\sqrt{\Sigma(p_i - q_i)^2}$. This was labeled 'TV' throughout all output documents. The error was identified during a ChatGPT corpus review on March 22, 2026.

Why it matters

Total variation distance for probability distributions is $\frac{1}{2} \Sigma|p_i - q_i|$ — the half-L1 norm, not the L2 norm. These are different metrics with different mathematical properties:

Metric	Definition and behavior
L2 drift (Euclidean)	$\sqrt{\Sigma(p_i - q_i)^2}$. Amplifies large single-component changes (squaring effect). More sensitive to concentrated shifts.
TV distance (half-L1)	$\frac{1}{2} \Sigma\|p_i - q_i\|$. Treats all component changes linearly. Bounded [0, 1] for probability vectors. Standard in information theory.

Mathematical relationship: for probability vectors, $TV \leq L2 \cdot \sqrt{n/2}$ and $L2 \leq \sqrt{2} \cdot TV$. They are equivalent up to dimension-dependent constants but not identical.

The correction

- All instances of the 'tv' metric renamed to 'l2_drift'.
- New tv_distance metric computed as $\frac{1}{2} \Sigma|p_i - q_i|$ added to all scripts.
- Both metrics now computed side-by-side in all analysis scripts (EMBER v3, OWID v2, GDP v2).
- All output documents regenerated with corrected naming.

Impact on results

The two metrics agree on all shock hit/miss verdicts across all three domains tested (EMBER, OWID, GDP). Drift counts differ in a few countries where TVD picks up additional drift years. No previously significant p-value is invalidated. One new divergence found in China OWID electricity data.

Key finding: the mislabeling did not invalidate results, but it could have undermined credibility with any reviewer who noticed that the 'TV' function was computing L2. The correction is presented here as evidence of methodological self-discipline, not as a minor fix.

Why this matters for CoDaWork

For a CoDa audience, showing that the project caught its own metric-naming error, documented both metrics side-by-side, and regenerated all outputs is more persuasive than trying to look finished. It says the project can survive methodological criticism — which is exactly what this packet is asking for.

Full corpus, analyzer tool, and documented failure modes: github.com/PeterHiggins19/Higgins-Unity-Framework

Peter Higgins | Rogue Wave Audio | PeterHiggins@RogueWaveAudio.com | March 2026